

CSCI4156 Assignment 1 – Canonical Multi-layer Perceptron

Instructor: Malcolm Heywood (mheywood@cs.dal.ca)

Due: 16 May – BrightSpace

Build your own Multi-Layer Perceptron

The purpose of this assignment is to provide a canonical neural machine learning model that we will incrementally adapt to support reinforcement learning. The pre-requisite class (csci3151) has already established the following for a multi-layer perceptron (MLP):

- **Representation:** two layers of ‘perceptrons’ where each perceptron performs the following mapping from a vector of inputs (\vec{x}) to scalar output or $y = f(net)$ where $net = \vec{w} \cdot \vec{x} + w_0$ and $\langle \vec{w}, w_0 \rangle$ denote the learning parameters (weights) that will be adjusted during credit assignment.
 - The number of inputs at the first layer is defined a priori by the task domain.
 - Likewise, the number of outputs is also set by the task domain.
 - The transfer (or activation) function $f(\cdot)$ takes the form of the sigmoid operator, or $f(net) = (1 + \exp(-net))^{-1}$
- **Cost function:** defines the performance criterion as empirical error minimization, or the specific case of $J_{SSE} = \frac{1}{2} \sum_{p \in N} e_p^2$
- **Credit Assignment:** is the process by which learning parameters $\langle \vec{w}, w_0 \rangle$ are adapted. Such a process can be considered a form of optimization (given the the model has a fixed number of learning parameters).
 - We are interested in solving this iteratively using the (stochastic) gradient decent heuristic.
 - This implies that for each exemplar a back-propagation weight update will be performed.

The Ask

The objective of assignment 1 is to develop your own software that implements the multi-layer perceptron with a single hidden layer given the following arguments:

1. number of input features
2. number of outputs
3. number of hidden layer neurons
4. maximum number of passes through the training partition

5. training dataset name as a .csv file
6. test dataset name as a .csv file
7. if no arguments are supplied, you should specify the argument sequence and meaning (which should be as above)

You should assume a Sigmoid transfer (activation) function and implement your MLP in Python **without** the use of any machine learning specific libraries (e.g. TensorFlow or Keras). Your submission will be evaluated on the Iris classification dataset¹ and a second mystery dataset. You should assume that class labels are encoded using a one hot encoding (e.g. for a three class problem $\vec{y} = \{\{1, 0, 0\}, \{0, 1, 0\}, \{0, 0, 1\}, \}$).

It is recommended that training data is subject to a feature-wise z-score normalization, or

$$x'_{i,p} = \frac{x_{i,p} - \mu_i}{\sigma_i} \quad (1)$$

where i is the feature and p is the data instance. μ and σ are the feature specific mean and variance as estimated over the data partition.

Post training, you should be able to take a second .csv file (the test partition) and suggest labels for each data record. Evaluation of your software will be performed under the undergraduate server (timberlea.cs.dal.ca). In particular, you should not assume any support for third party libraries (e.g. TensorFlow or Keras). Attempting to do so will result in a zero grade.

Your predicted labels for the test file should appear in an output .csv file with the file name: “B00XXXXXX” where the ‘X’ values correspond to your banner number.

¹Included with assignment as a .csv file. First 4 attributes are input features, last 3 attributes are output labels.